

Gur Raunaq Singh

+91-9015154091 | raunaq.soni@gmail.com | linkedin.com/in/gurraunaqsingh | github.com/raunaqness | Bangalore, India

SUMMARY

Senior AI and ML Systems Architect with 8+ years of experience designing secure, scalable AI platforms in enterprise environments processing high-volume transactional data. Specialized in distributed ML systems, Kubernetes-native orchestration, multi-cloud architectures, and Generative AI (LLMs, RAG, agentic workflows). Strong focus on MLOps, IAM-based access control, auditability, and operational resilience. Experience deploying and managing containerized workflow automation systems (n8n) on cloud infrastructure with secure API and LLM integrations.

TECHNICAL SKILLS

AI/ML: LLMs, RAG, LangChain, LangGraph, Model Monitoring

Cloud & Infra: Azure (AKS, App Services), AWS, Kubernetes, Docker, Terraform

Data & Backend: Python, SQL, Kafka, Redis, PostgreSQL, FastAPI, Django

MLOps: CI/CD, MLflow, IAM, Audit Logging, Encryption, SLA Monitoring

Automation: n8n (workflow automation), API integrations, Webhooks

EXPERIENCE

Aptos India

Machine Learning Engineer

Jan 2023 – Present

Bangalore, India

Project: Aptos One – Enterprise Unified Commerce Platform serving 100+ global retailers and processing billions of transactions annually.

- Architected a Kubernetes-native ML orchestration control plane enabling workload-aware scheduling, checkpoint-based recovery, and distributed job lifecycle management in SLA-bound production environments.
- Designed secure, fault-tolerant data pipelines on GCP (BigQuery + event-driven workflows), implementing IAM-based access controls, encryption practices, and audit logging while improving processing SLAs by 35%.
- Optimized cloud infrastructure and query execution strategies, reducing compute costs by 60% (about USD 240K annual savings) while modernizing legacy systems into scalable, observable cloud-native components.
- Technologies: Python, GCP, Vertex AI, BigQuery, Kubernetes, Docker, Terraform, Redis, CI/CD, IAM

Project: Generative AI & LLM Support Assistant

- Designed an enterprise-grade RAG architecture integrating structured analytics metadata and internal knowledge bases to deliver context-aware responses with improved semantic accuracy.
- Implemented prompt versioning, evaluation workflows, and monitoring mechanisms to reduce hallucinations by 35% while improving traceability and observability of model outputs.
- Deployed containerized inference workflows using MLOps practices, enabling controlled experimentation, scalable inference, and secure knowledge isolation via vector database access controls.
- Technologies: Python, LangChain, LangGraph, Vertex AI, VectorDB, FastAPI, Docker, MLOps

Turing.com

Senior Software Engineer

Apr 2022 – Dec 2022

India (Remote)

Project: End-to-End ML Deployment Platform

- Designed and implemented production-grade ML pipelines using Airflow for ingestion, training, validation, and automated deployment.
- Standardized CI/CD and model versioning practices, reducing deployment time by 60% and improving controlled production rollouts.
- Integrated monitoring and validation checks to enhance production reliability and maintainability of deployed ML systems.
- Technologies: Python, Django, Airflow, GCP, Docker, MLOps

Clickpost Pvt. Ltd.

Senior Software Engineer

Apr 2021 – Apr 2022

Bangalore, India

Project: Cash Reconciliation Automation Platform – serving 1,000+ enterprise customers handling high-volume CoD transactions.

- Designed and built a high-accuracy financial reconciliation system processing large-scale transactional datasets with strict data integrity and traceability requirements.
- Reduced reconciliation cycle time by 95% (from about 2 days to minutes) through parallelized processing and automated validation workflows, improving operational efficiency and financial visibility.
- Architected secure ingestion pipelines and scalable backend APIs under high-availability constraints, ensuring auditability and controlled access to financial data.
- Technologies: Python, Django, PostgreSQL, AWS, Redshift, Kafka, Docker, Kubernetes

Veda Labs Pvt. Ltd.

Software Engineer

Jul 2018 – Apr 2021

New Delhi, India

Project: Video Analytics Platform – Computer vision platform deployed on-premise for retail analytics using CCTV footage.

- Developed containerized video analytics pipelines for edge devices.
- Built a Kafka-based streaming architecture for scalable, real-time processing.
- Technologies: Python, OpenCV, Computer Vision, Kafka, Docker

PROJECTS

Stateful AI Travel Agent (LangGraph)

- Engineered a graph-orchestrated, memory-aware multi-agent system for dynamic travel planning.
- Github: <https://github.com/raunaqness/ai-travel-agent>
- Demo: <https://ai-travel-agent.fly.dev/>

Workflow Automation Platform (n8n + Azure)

- Designed and deployed containerized n8n workflows on Azure integrating APIs and LLM services for secure, scalable process automation.

EDUCATION

Indraprastha University

Bachelor of Technology – Computer Science Engineering

Aug 2014 – Jun 2018

New Delhi, India